

Package ‘sentometrics’

October 26, 2017

Title An Integrated Framework for Textual Sentiment Time Series Aggregation and Forecasting

Version 0.2.0

Date 2017-11-01

Description Time series analysis based on textual sentiment, accounting for the intrinsic challenge that sentiment can be computed and pooled across texts and time in many ways. Provides a means to model the impact of sentiment in texts on a target variable, by first computing a wide range of textual sentiment measures and then selecting those that are most informative.

Depends R (>= 3.4.2), data.table, ggplot2, foreach

License GPL-2

Encoding UTF-8

LazyData true

Suggests testthat, e1071, randomForest

Imports utils, stats, quanteda, sentimentr, stringi, zoo, abind, glmnet, caret, compiler, Rcpp (>= 0.12.13), RcppRoll, ggthemes, ISOweek, MCS

LinkingTo Rcpp, RcppArmadillo

RoxygenNote 6.0.1

R topics documented:

add_features	2
almons	3
compute_sentiment	4
ctr_agg	5
ctr_merge	7
ctr_model	8
e1071	10
exponentials	11
fill_measures	11
get_hows	12
lexicons	13
merge_measures	14
perform_agg	15
perform_MCS	16
plot.sentomeasures	16
plot.sentomodeliter	17
plot_attributions	19

predict.sentomodel	19
retrieve_attributions	20
scale.sentomeasures	21
select_measures	22
sentometrics	23
sento_corpus	24
sento_measures	25
sento_model	27
setup_lexicons	30
sp500	31
to_global	32
usnews	33
valence	34

Index	35
--------------	-----------

add_features	<i>Add feature columns to a sentocorpus</i>
--------------	---

Description

Adds new named feature columns to provided sentocorpus object.

Usage

```
add_features(sentocorpus, featuresdf = NULL, keywords = NULL)
```

Arguments

sentocorpus	a sentocorpus object.
featuresdf	a named data.frame with as columns the new features of type numeric to add to the sentocorpus inputted. If the number of rows in featuresdf is not equal to the number of documents in sentocorpus, recycling will occur.
keywords	a named character vector. For every element in this vector, a new feature column will be added with as value 1 for the texts in which the element (e.g. a word) appears, and 0 if not. The names of the vector are used as the names of the new features.

Details

If a provided feature name is already part of the corpus, it will be replaced. The featuresdf and keywords arguments can be provided at the same time, or only one of them, leaving the other at NULL.

Value

An updated sentocorpus object.

Author(s)

Samuel Borms

Examples

```
data("usnews")

# construct a corpus and add random features to it
corpus <- sento_corpus(corpusdf = usnews)
corpus <- add_features(corpus,
  featuresdf = data.table(random = runif(quanteda::ndoc(corpus))))
corpus <- add_features(corpus,
  featuresdf = data.table(random = runif(quanteda::ndoc(corpus))),
  keywords = c(pres = "president", war = "war"))
```

almons

Compute Almon polynomials

Description

Computes Almon polynomial weighting curves; handy to self-select specific time aggregation weighting schemes.

Usage

```
almons(n, orders = 1:3, do.inverse = TRUE, do.normalize = TRUE)
```

Arguments

n	a single numeric to indicate the length of the curve (the number of lags, cf. n in the formula).
orders	a numeric vector as the sequence the Almon orders (cf. b in the formula).
do.inverse	TRUE if the inverse Almon polynomials should be calculated as well.
do.normalize	TRUE if polynomials should be normalized to unity.

Details

The Almon polynomial formula implemented is: $(1 - (i/n)^b)(i/n)^{max(b)-b}$.

Value

A data.frame of all Almon polynomial weighting curves, of size length(orders) (times two if do.inverse == TRUE).

See Also

[ctr_agg](#)

compute_sentiment	<i>Compute document-level sentiment across features and lexicons</i>
-------------------	--

Description

Given a corpus of texts, computes sentiment per document using the bag-of-words approach, based on the lexicons provided and a choice of aggregation across words per document scheme. Relies partly on the **quanteda** package. The scores computed are net sentiment (sum of positive minus sum of negative scores).

Usage

```
compute_sentiment(sentocorpus, lexicons, how = get_hows()$words, dfm = NULL)
```

Arguments

sentocorpus	a sentocorpus object.
lexicons	output from a <code>setup_lexicons()</code> call.
how	a single character vector defining how aggregation within documents should be performed. For currently available options on how aggregation can occur, see <code>get_hows()\$words</code> .
dfm	optional; an output from a <code>quanteda::dfm()</code> call, such that users can specify their own tokenization scheme (via <code>quanteda::tokenize()</code>) as well as other parameters related to the construction of a document-feature matrix (dfm). By default, a dfm is created based on a tokenization that removes punctuation, numbers, symbols and separators. We suggest to stick to unigrams, as the remainder of the sentiment computation and built-in lexicons assume the same.

Details

For a separate calculation of positive (resp. negative) sentiment, one has to provide distinct positive (resp. negative) lexicons. This can be done using the `do.split` option in the `setup_lexicons()` function, which splits out the lexicons into a positive and a negative polarity counterpart. NAs are converted to 0, under the assumption that this is equivalent to no sentiment.

Value

A list containing:

corpus	the supplied sentocorpus object; the texts are altered if valence shifters are part of the lexicons.
sentiment	the sentiment scores data.table with a date and lexicon-feature sentiment scores columns.
features	a character vector of the different features.
lexicons	a character vector of the different lexicons used.
howWithin	a character vector to remind how sentiment within documents was aggregated.

Author(s)

Samuel Borms

See Also[dfm](#), [tokenize](#)**Examples**

```
data("usnews")
data("lexicons")
data("valence")

# sentiment computation based on raw frequency counts
corpus <- sento_corpus(corpusdf = usnews)
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
sent <- compute_sentiment(corpus, l, how = "counts")

# same sentiment computation based on a user-supplied dfm with default settings
dfm <- quanteda::dfm(quanteda::tokenize(corpus), verbose = FALSE)
sent <- compute_sentiment(corpus, l, how = "counts", dfm = dfm)
```

ctr_agg

*Set up control for aggregation into sentiment measures***Description**

Sets up control object for aggregation of document-level textual sentiment into textual sentiment measures (indices).

Usage

```
ctr_agg(howWithin = "proportional", howDocs = "equal_weight",
  howTime = "equal_weight", do.ignoreZeros = FALSE, by = "day", lag = 1,
  fill = "zero", alphasExp = seq(0.1, 0.5, by = 0.1), ordersAlm = 1:3,
  do.inverseAlm = TRUE, do.normalizeAlm = TRUE, weights = NULL,
  dfm = NULL)
```

Arguments

- | | |
|----------------|---|
| howWithin | a single character vector defining how aggregation within documents will be performed. Should <code>length(howWithin) > 1</code> , the first element is used. For currently available options on how aggregation can occur, see <code>get_hows()\$words</code> . |
| howDocs | a single character vector defining how aggregation across documents per date will be performed. Should <code>length(howDocs) > 1</code> , the first element is used. For currently available options on how aggregation can occur, see <code>get_hows()\$docs</code> . |
| howTime | a character vector defining how aggregation across dates will be performed. More than one choice is possible here. For currently available options on how aggregation can occur, see <code>get_hows()\$time</code> . |
| do.ignoreZeros | a logical indicating whether zero sentiment values have to be ignored in the determination of the document weights while aggregating across documents. |
| by | a single character vector, either "day", "week", "month" or "year", to indicate at what level the dates should be aggregated. Dates will be displayed as the first day of the period, if applicable (e.g. "2017-03-01" for March 2017). |

lag	a single integer vector, being the time lag to be specified for aggregation across time. By default equal to 1, meaning no aggregation across time.
fill	a single character vector, one of <code>c("zero", "latest", "none")</code> , to control how missing sentiment values across the continuum of dates considered are added. This impacts the aggregation across time, applying the <code>fill_measures()</code> function before aggregating, except if <code>fill == "none"</code> . By default equal to "zero", which sets the scores (and thus also the weights) of the added dates to zero in the time aggregation.
alphasExp	a numeric vector of all exponential smoothing factors to calculate weights for, used if "exponential" %in% <code>howTime</code> . Values should be between 0 and 1 (both excluded).
ordersAlm	a numeric vector of all Almon polynomial orders to calculate weights for, used if "almon" %in% <code>howTime</code> .
do.inverseAlm	a logical indicating if for every Almon polynomial its inverse has to be calculated too, used if "almon" %in% <code>howTime</code> .
do.normalizeAlm	a logical indicating if every Almon polynomial weights column should sum to one, used if "almon" %in% <code>howTime</code> .
weights	an optional own weighting scheme, always used if provided as a <code>data.frame</code> with the number of rows equal to the desired lag. The automatic Almon polynomials and exponential weighting functions are created sequentially; if the user wants only specific of such time weighting series it can use <code>almons()</code> and <code>exponentials()</code> , select the columns it requires, combine it into a <code>data.frame</code> and supply it under this argument (see examples).
dfm	optional; an output from a <code>quanteda::dfm()</code> call, such that users can specify their own tokenization scheme (via <code>quanteda::tokenize()</code>) as well as other parameters related to the construction of a document-feature matrix (dfm). By default, a dfm is created based on a tokenization that removes punctuation, numbers, symbols and separators. We suggest to stick to unigrams, as the remainder of the sentiment computation (in <code>compute_sentiment()</code>) and built-in lexicons assume the same.

Details

For currently available options on how aggregation can occur (via the `howWithin`, `howDocs` and `howTime` parameters), call `get_hows()`.

Value

A list encapsulating the control parameters.

Author(s)

Samuel Borms, Keven Bluteau

See Also

[fill_measures](#), [compute_sentiment](#)

Examples

```
# simple control function
ctr1 <- ctr_agg(howTime = "linear", by = "year", lag = 3)

# more elaborate control function (particular attention to time weighting schemes)
ctr2 <- ctr_agg(howWithin = "tf-idf",
               howDocs = "proportional",
               howTime = c("equal_weight", "linear", "almon", "exponential", "own"),
               do.ignoreZeros = TRUE,
               by = "day",
               lag = 20,
               ordersAlm = 1:3,
               do.inverseAlm = TRUE,
               do.normalizeAlm = TRUE,
               alphasExp = c(0.20, 0.50, 0.70, 0.95),
               weights = data.frame(myWeights = runif(20)))

# set up control function with one linear, two Almon and two exponential weighting schemes
a <- almons(n = 70, orders = 1:3, do.inverse = TRUE, do.normalize = TRUE)
e <- exponentials(n = 70, alphas = c(0.4, 0.8))
ctr3 <- ctr_agg(howTime = c("linear"), by = "year",
               weights = data.frame(a1 = a[, 1], a2 = a[, 3], e1 = e[, 1], e2 = e[, 2]))
```

ctr_merge

Set up control for merging sentiment measures

Description

Sets up control object for the optional merging (additional aggregation) of sentiment measures.

Usage

```
ctr_merge(sentomeasures, feat = NA, lex = NA, time = NA,
          do.keep = FALSE)
```

Arguments

sentomeasures	a sentomeasures object. This is necessary to check whether the other input arguments make sense.
feat	a list with unique features to merge at given name, e.g. <code>list(feat12 = c("feat1", "feat2"))</code> . See <code>sentomeasures\$features</code> for the exact names to use.
lex	a list with unique lexicons to merge at given name, e.g. <code>list(lex12 = c("lex1", "lex2"))</code> . See <code>sentomeasures\$lexicons</code> for the exact names to use.
time	a list with unique time weighting schemes to merge at given name, e.g. <code>list(tw12 = c("tw1", "tw2"))</code> . See <code>sentomeasures\$time</code> for the exact names to use.
do.keep	a logical indicating if the original sentiment measures should be kept (i.e. the merged sentiment measures will be added to the current sentiment measures as additional indices if TRUE).

Value

A list encapsulating the control parameters.

Author(s)

Samuel Borms

See Also

[merge_measures](#)

Examples

```
data("usnews")
data("lexicons")
data("valence")

# construct a sentomeasures object to start with
corpus <- sento_corpus(corpusdf = usnews)
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
ctr <- ctr_agg(howTime = c("equal_weight", "linear"), by = "year", lag = 3)
sentomeasures <- sento_measures(corpus, l, ctr)

# set up a correct control function
ctrMerge <- ctr_merge(sentomeasures,
                      time = list(W = c("equal_weight", "linear")),
                      lex = list(LEX = c("LM_eng", "HENRY_eng")),
                      feat = list(journals = c("wsj", "wapo")),
                      do.keep = TRUE)

## Not run:
# produces an informative error message
ctrMerge <- ctr_merge(sentomeasures,
                      time = list(W = c("equal_weight", "almon1")),
                      lex = list(LEX = c("LM_eng", "HENRY_eng")),
                      feat = list(journals = c("notInHere", "wapo")))

## End(Not run)
```

ctr_model

Set up control for sentiment measures-based regression modelling

Description

Sets up control object for linear or nonlinear modelling of a response variable onto a large panel of textual sentiment measures (and potentially other variables). See the [sento_model](#) function's documentation for details on the estimation and calibration procedure.

Usage

```
ctr_model(model = c("gaussian", "binomial", "multinomial"), type = c("BIC",
  "AIC", "Cp", "cv"), intercept = TRUE, do.iter = FALSE, h = 0,
  alphas = seq(0, 1, by = 0.2), nSample = NULL, trainWindow = NULL,
  testWindow = NULL, oos = 0, start = 1, do.progress = TRUE,
  do.parallel = FALSE)
```

Arguments

model	a character vector with one of the following: "gaussian" (linear regression), "binomial" (binomial logistic regression), or "multinomial" (multinomial logistic regression).
type	a character vector indicating which model selection criteria to use. Currently supports "BIC", "AIC" and "Cp" (Mallows's Cp) as sparse regression adapted information criteria (cf. "On the 'degrees of freedom' of the LASSO"; Zou, Hastie, Tibshirani et al., 2007), and "cv" (cross-validation based on the train function from the caret package). The adapted information criteria are currently only available for a linear regression.
intercept	a logical, TRUE by default fits an intercept.
do.iter	a logical, TRUE induces an iterative estimation of models at the given nSample size and performs the associated one-step ahead out-of-sample prediction exercise through time.
h	an integer value to shift the time series to have the desired prediction setup, $h == 0$ means no change to the input data (nowcasting assuming data is aligned properly), $h > 0$ shifts the dependent variable by h periods (i.e. rows) further in time (forecasting), $h < 0$ shifts the independent variables by h periods.
alphas	a numeric vector of the different alphas to test for during optimization, between 0 and 1. A value of 0 pertains to Ridge optimization, a value of 1 to LASSO optimization; values in between are pure elastic net. The lambda values tested for are automatically chosen by the <code>glmnet()</code> function or set to $10^{\text{seq}(2, -2, \text{length.out} = 100)}$ in case of cross-validation.
nSample	a positive integer as the size of the sample for model calibration at every iteration (ignored if <code>iter == FALSE</code>).
trainWindow	a positive integer as the size of the training sample in cross-validation (ignored if <code>type != "cv"</code>).
testWindow	a positive integer as the size of the test sample in cross-validation (ignored if <code>type != "cv"</code>).
oos	a non-negative integer to indicate the number of periods to skip from the end of the cross-validation training sample (out-of-sample) up to the test sample (ignored if <code>type != "cv"</code>).
start	a positive integer to indicate at which point the iteration has to start (ignored if <code>iter == FALSE</code>). For example, for 100 out-of-sample iterations, <code>start = 70</code> only performs the analysis for the last 31 samples.
do.progress	a logical, if TRUE progress statements are displayed during model calibration.
do.parallel	a logical, if TRUE the <code>%dopar%</code> construct from the foreach package is applied for iterative model estimation. A proper parallel backend needs to be set up to make it work. No progress statements are displayed whatsoever when TRUE. For cross-validation models, parallelization is also carried out for single-run models, whenever a parallel backend is set up. See the examples in sento_model .

Value

A list encapsulating the control parameters.

Author(s)

Samuel Borms, Keven Bluteau

See Also

[sento_model](#)

Examples

```
# series of example information criterion based model control functions
ctrIC1 <- ctr_model(model = "gaussian", type = "BIC", do.iter = FALSE, h = 0,
                    alphas = seq(0, 1, by = 0.10))
ctrIC2 <- ctr_model(model = "gaussian", type = "BIC", do.iter = TRUE, h = 0, nSample = 100)

# series of example cross-validation based model control functions
ctrCV1 <- ctr_model(model = "gaussian", type = "cv", do.iter = FALSE, h = 0, trainWindow = 250,
                    testWindow = 4, oos = 0, do.progress = TRUE)
ctrCV2 <- ctr_model(model = "binomial", type = "cv", h = 0, trainWindow = 250,
                    testWindow = 4, oos = 0, do.progress = TRUE)
ctrCV3 <- ctr_model(model = "multinomial", type = "cv", h = 0, trainWindow = 250,
                    testWindow = 4, oos = 0, do.progress = TRUE)
ctrCV4 <- ctr_model(model = "gaussian", type = "cv", do.iter = TRUE, h = 0, trainWindow = 45,
                    testWindow = 4, oos = 0, nSample = 70, do.progress = TRUE)
```

epu

Monthly Economic Policy Uncertainty Index

Description

Monthly values of a news-based index of US Economic Policy Uncertainty (EPU) between January 1980 and September 2014, including a binomial and a multinomial example series. For more information on its calculation, see [this](#). Following columns are present:

- date. Date as "yyyy-mm-01".
- index A numeric monthly index value.
- above. A factor with value "above" if the index is greater than the mean of the entire series, else "below".
- aboveMulti. A factor with values "above+", "above", "below" and "below-" if the index is greater than the 75 respectively and in a mutually exclusive sense.

Usage

```
data("epu")
```

Format

A data.frame with 417 rows and 4 columns.

Source

Research on Economic Policy Uncertainty

exponentials	<i>Compute exponential weighting curves</i>
--------------	---

Description

Computes exponential weighting curves; handy to self-select specific time aggregation weighting schemes.

Usage

```
exponentials(n, alphas = seq(0.1, 0.5, by = 0.1))
```

Arguments

`n` a single numeric to indicate the length of the curve (the number of lags).
`alphas` a numeric vector of decay factors.

Value

A data.frame of exponential weighting curves per value of alphas.

See Also

[ctr_agg](#)

fill_measures	<i>Add and fill missing dates</i>
---------------	-----------------------------------

Description

Adds missing dates between earliest and latest date, such that time series is continuous on a period-by-period basis. Fills in these dates with either 0, the respective latest non-missing value or NA.

Usage

```
fill_measures(sentomeasures, fill = "zero")
```

Arguments

`sentomeasures` a sentomeasures object.
`fill` an element of `c("zero", "latest", NA)`; the first and last assume missing dates represent zero sentiment, the second assumes missing dates represent constant sentiment.

Value

A modified sentomeasures object.

Author(s)

Samuel Borms

Examples

```
data("usnews")
data("lexicons")
data("valence")

# construct a sentomeasures object to start with
corpus <- sento_corpus(corpusdf = usnews)
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
ctr <- ctr_agg(howTime = c("equal_weight", "linear"), by = "year", lag = 3)
sentomeasures <- sento_measures(corpus, l, ctr)

# fill measures
f1 <- fill_measures(sentomeasures)
f2 <- fill_measures(sentomeasures, fill = "latest")
f3 <- fill_measures(sentomeasures, fill = NA)
```

get_hows

Options supported to perform aggregation into sentiment measures.

Description

Call for information purposes only. Used within `ctr_agg()` to check if supplied aggregation hows are supported.

Usage

```
get_hows()
```

Value

A list with the supported aggregation hows for arguments `howWithin` (within documents), `howDows` (across documents, per date) and `howTime` (across dates), to be supplied to `ctr_agg()`.

See Also

[ctr_agg](#)

`lexicons`*Built-in lexicons*

Description

A list containing all built-in lexicons as a `data.table` with two columns: a `x` column with the words, and a `y` column with the polarities. The list element names incorporate consecutively the name and language, and `"_tr"` as suffix if the lexicon is translated. The lexicons are in the form required for further sentiment analysis. The built-in lexicons are the following:

- `FEEL_eng_tr` (FEEL: French Expanded Emotion Lexicon)
- `FEEL_fr`
- `FEEL_nl_tr`
- `GI_eng` (GI: General Inquirer, i.e. Harvard IV-4 combined with Laswell)
- `GI_fr_tr`
- `GI_nl_tr`
- `HENRY_eng` (HENRY: Henry)
- `HENRY_fr_tr`
- `HENRY_nl_tr`
- `LM_eng` (LM: Loughran and McDonald)
- `LM_fr_tr`
- `LM_nl_tr`

Usage

```
data("lexicons")
```

Format

A list with all built-in lexicons, appropriately named as `"NAME_language(_tr)"`.

Source

[FEEL lexicon](#)

[GI lexicon](#)

[HENRY lexicon](#)

[LM lexicon](#)

Examples

```
lexicons[c("FEEL_eng_tr", "LM_eng")]
```

merge_measures

Merge sentiment measures

Description

Merge (further aggregate) measures by combining across the lexicons, features and time weighting schemes dimensions. The combination occurs by taking the mean of the relevant measures.

Usage

```
merge_measures(ctr)
```

Arguments

ctr output from a ctr_merge() call.

Value

A modified sentomeasures object, with only the sentiment measures required, including updated information and statistics, but the original sentiment scores data.table untouched.

Author(s)

Samuel Borms

See Also

[ctr_merge](#)

Examples

```
data("usnews")
data("lexicons")
data("valence")

# construct a sentomeasures object to start with
corpus <- sento_corpus(corpusdf = usnews)
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
ctr <- ctr_agg(howTime = c("equal_weight", "linear"), by = "year", lag = 3)
sentomeasures <- sento_measures(corpus, l, ctr)

# set up control function and perform the merging
ctrMerge <- ctr_merge(sentomeasures,
  time = list(W = c("equal_weight", "linear")),
  feat = list(journals = c("wsj", "wapo")),
  do.keep = TRUE)
sentomeasuresMerged <- merge_measures(ctrMerge)
```

`perform_agg`*Aggregate textual sentiment across documents and time*

Description

Condense document-level textual sentiment scores into a panel of textual sentiment measures by aggregating across documents and time.

Usage

```
perform_agg(toAgg, ctr)
```

Arguments

<code>toAgg</code>	output from a <code>compute_sentiment()</code> call, a list with as main component a sentiment scores <code>data.table</code> with dates and lexicon-feature sentiment scores columns.
<code>ctr</code>	output from a <code>ctr_agg()</code> call.

Value

A `sentomeasures` object.

Author(s)

Samuel Borms, Keven Bluteau

See Also

[compute_sentiment](#), [ctr_agg](#), [sento_measures](#)

Examples

```
data("usnews")
data("lexicons")
data("valence")

# computation of sentiment and aggregation into sentiment measures
corpus <- sento_corpus(corpusdf = usnews)
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
sent <- compute_sentiment(corpus, l, how = "counts")
ctr <- ctr_agg(howTime = c("linear"), by = "year", lag = 3)
sentomeasures <- perform_agg(sent, ctr)
```

perform_MCS	<i>Apply the model confidence set (MCS) procedure to a selection of models</i>
-------------	--

Description

Calculates the model confidence set (see “The Model Confidence Set”; Hansen, Lunde and Nason, 2011) as implemented in the **MCS** package, for a set of different sentomodeliter objects.

Usage

```
perform_MCS(models, loss = c("DA", "errorSq", "AD", "accuracy"), ...)
```

Arguments

models	a named list of sentomodeliter objects. All models should be of the same family, being either "gaussian", "binomial" or "multinomial", and have performance data of the same dimensions.
loss	a single character vector, either "DA" (directional <i>inaccuracy</i>), "errorSq" (squared errors), "AD" (absolute errors) or "accuracy" (<i>inaccurate</i> class predictions). This argument defines on what basis the model confidence set is calculated. The first three options are available for "gaussian" models, the last option applies only to "binomial" or "multinomial" models.
...	other parameters that can be supplied to the MCS::MCSprocedure() function.

Value

An object as returned from the MCS::MCSprocedure() function.

Author(s)

Samuel Borms

See Also

[sento_model](#), [MCSprocedure](#)

plot.sentomeasures	<i>Plot sentiment measures</i>
--------------------	--------------------------------

Description

Straightforward plotting method that shows all sentiment measures from the provided sentomeasures object in one plot, or the average along one of the lexicons, features and time weighting dimensions. We suggest to make use of the codeselect_measures() function when you desire to plot only a subset of the sentiment measures.

Usage

```
## S3 method for class 'sentomeasures'
plot(x, group = "all", ...)
```


Arguments

x	a sentomeasures object.
group	a value from c("lexicons", "features", "time", "all"). The first three choices display the average of all measures from the same group, in a different color. The choice "all" displays every single sentiment measure in a separate color, but this may look visually overwhelming very fast, and can be quite slow.
...	not used.

Value

Returns a simple **ggplot2** plot, which can be added onto (or to alter its default elements) by using the + operator (see examples). By default, a legend is positioned at the top if there are at maximum twelve line graphs plotted.

Author(s)

Samuel Borms

Examples

```
data("usnews")
data("lexicons")
data("valence")

# construct a sentomeasures object to start with
corpus <- sento_corpus(corpusdf = usnews)
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
ctr <- ctr_agg(howTime = c("equal_weight", "linear"), by = "year", lag = 3)
sentomeasures <- sento_measures(corpus, l, ctr)

# plot sentiment measures
plot(sentomeasures)
plot(sentomeasures, group = "lexicons")

# adjust appearance of plot
p <- plot(sentomeasures)
p <- p +
  ggthemes::theme_base() +
  scale_x_date(name = "date") +
  scale_y_continuous(name = "newName")
p
```

plot.sentodeliter *Plot iterative predictions versus realized values*

Description

Displays a plot of all predictions made through the iterative model computation as incorporated in the input sentodeliter object, as well as the corresponding true values.

Usage

```
## S3 method for class 'sentomodeliter'
plot(x, ...)
```

Arguments

```
x                a sentomeasures object.
...              not used.
```

Value

Returns a simple **ggplot2** plot, which can be added onto (or to alter its default elements) by using the + operator (see examples).

Author(s)

Samuel Borms

Examples

```
data("usnews")
data("lexicons")
data("valence")
data("epu")

# construct a sentomeasures object to start with
corpusAll <- sento_corpus(corpusdf = usnews)
corpus <- quantda::corpus_subset(corpusAll, date >= "1980-01-01" & date < "2014-10-01")
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
ctr <- ctr_agg(howWithin = "tf-idf", howDocs = "proportional",
               howTime = c("equal_weight", "linear", "almon"),
               by = "month", lag = 3, ordersAlm = 1:3,
               do.inverseAlm = TRUE, do.normalizeAlm = TRUE)
sentomeasures <- sento_measures(corpus, l, ctr)

# prepare y and other x variables
y <- epu[epu$date >= sentomeasures$measures$date[1], ]$index
length(y) == nrow(sentomeasures$measures) # TRUE
x <- data.frame(runif(length(y)), rnorm(length(y))) # two other (random) x variables
colnames(x) <- c("x1", "x2")

# estimate regression iteratively based on a sample of 120, skipping the first 275 samples
ctr <- ctr_model(model = "gaussian", type = "AIC", do.iter = TRUE,
                 h = 0, nSample = 120, start = 276)
out <- sento_model(sentomeasures, y, x = x, ctr = ctr)
summary(out)

# plotting
p <- plot(out)
p <- p +
  ggthemes::theme_few()
p
```

plot_attributions	<i>Plot prediction attribution at specified level</i>
-------------------	---

Description

Shows a plot of the attributions along the dimension provided.

Usage

```
plot_attributions(attributions, group = "features")
```

Arguments

attributions	an output from a <code>retrieve_attributions()</code> call.
group	a value from <code>c("lexicons", "features", "time", "lags")</code> .

Details

See [sento_model](#) for an elaborate modelling example including the calculation and plotting of attributions. This function does not handle the plotting of the attribution of individual documents, since there are often a lot of documents involved and de facto they appear only once at one date (even though a document may contribute to predictions at several dates, depending on the number of lags in the time aggregation).

Value

Returns a simple **ggplot2** plot, which can be added onto (or to alter its default elements) by using the `+` operator (see examples). By default, a legend is positioned at the top if the number of dimensions (thus, individual plots) is at maximum twelve.

Author(s)

Samuel Borms

<code>predict.sentomodel</code>	<i>Make predictions from a sentomodel object</i>
---------------------------------	--

Description

Prediction method for `sentomodel` class, with usage along the lines of `predict.glmnet`, but simplified in terms of allowed parameters.

Usage

```
## S3 method for class 'sentomodel'
predict(object, newx, type, offset = NULL, ...)
```

Arguments

<code>object</code>	a <code>sentomodel</code> object.
<code>newx</code>	a matrix of numeric values with all explanatory variables to be used for the prediction(s), structured row-by-row; see documentation for predict.glmnet . The number of variables should be equal to <code>sentomodel\$nVar</code> , being the sum of the number of original sentiment measures and the number of additional explanatory variables. Variables discarded in the regression process are discarded again here, based on <code>sentomodel\$discarded</code> .
<code>type</code>	type of prediction required, a value from <code>c("link", "response", "class")</code> , see documentation for predict.glmnet .
<code>offset</code>	not used. Any values here will be ignored.
<code>...</code>	not used.

Value

A prediction output depending on the `type` argument provided.

Author(s)

Samuel Borms

See Also

[predict.glmnet](#), [sento_model](#)

`retrieve_attributions` *Retrieve top-down sentiment attributions given model object*

Description

Computes the attributions to predictions for a (given) number of dates at all possible sentiment dimensions, based on the coefficients associated to each sentiment measure, as estimated in the provided model object.

Usage

```
retrieve_attributions(model, sentomeasures, do.normalize, refDates)
```

Arguments

<code>model</code>	a <code>sentomodel</code> or <code>sentomodeliter</code> object.
<code>sentomeasures</code>	the <code>sentomeasures</code> object used to estimate the <code>sentomodel</code> object argument.
<code>do.normalize</code>	a logical, TRUE divides each element of every attribution vector at a given date by its L2-norm at that date. The document attributions are not normalized.
<code>refDates</code>	the dates at which attribution is to be performed. These should be between the latest date available in the input <code>sentomeasures</code> object and the first date of the sample used to estimate the model, i.e. <code>model\$dates[1]</code> if <code>model</code> is a <code>sentomodel</code> object. All dates should also be present in <code>sentomeasures\$measures\$date</code> . If NULL (default), attribution is calculated for all in-sample dates. Ignored if <code>model</code> is a <code>sentomodeliter</code> object, for which attribution is calculated for all out-of-sample prediction dates.

Details

See [sento_model](#) for an elaborate modelling example including the calculation and plotting of attributions.

Value

A list with all possible dimensions for which aggregation is computed, being "documents", "lexicons", "features", "time" and "lags". The last four dimensions are `data.table`s having a "date" column and the other columns the different names of the dimension, with the attributions as values. For document-level attribution, the list is further decomposed into a `data.table` per date, with "id", date and attrib columns.

Author(s)

Samuel Borms, Keven Bluteau

See Also

[sento_model](#)

scale.sentomeasures	<i>Scaling and centering of sentiment measures</i>
---------------------	--

Description

Scales and centers the sentiment measures from a `sentomeasures` object, column-per-column. By default, the measures are normalized. NAs are removed first.

Usage

```
## S3 method for class 'sentomeasures'  
scale(x, center = TRUE, scale = TRUE)
```

Arguments

x	a <code>sentomeasures</code> object.
center	a logical, see documentation for the generic scale .
scale	a logical, see documentation for the generic scale .

Value

A modified `sentomeasures` object, with the measures replaced by the scaled measures as well as updated statistics.

Author(s)

Samuel Borms

Examples

```
data("usnews")
data("lexicons")
data("valence")

# construct a sentomeasures object to start with
corpus <- sento_corpus(corpusdf = usnews)
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
ctr <- ctr_agg(howTime = c("equal_weight", "linear"), by = "year", lag = 3)
sentomeasures <- sento_measures(corpus, l, ctr)

# scale sentimeant measures
scaled <- scale(sentomeasures)
```

select_measures	<i>Select a subset of sentiment measures</i>
-----------------	--

Description

Selects the subset of sentiment measures which include either all of the given selection components combined, or those who's name consist of at least one of the selection components. Selecting a subset of dates

Usage

```
select_measures(sentomeasures, toSelect = "all", do.combine = TRUE,
  dates = NA)
```

Arguments

sentomeasures	a sentomeasures object.
toSelect	a "character" vector of the lexicon, feature and time weighting scheme names, to indicate which measures need to be selected. By default equal to "all", which means no selection of the sentiment measures is made; this may be used if one only wants to extract a subset of dates via the dates argument.
do.combine	a logical indicating if only measures for wich all (TRUE) or at least one (FALSE) of the selection components should occur in each sentiment measure's name in the subset. If do.combine == TRUE, the toSelect argument can only consist of one lexicon, one feature and one time weighting scheme at maximum.
dates	any expression, in the form of a character vector, that would correctly evaluate to a logical vector, features the variable date and has dates specified as "yyyy-mm-dd", e.g. dates = "date >= '2000-01-15'". This argument may also be a vector of class Date which extracts all dates that show up in that vector. See the examples. By default equal to NA, meaning no subsetting based on dates is done.

Value

A modified sentomeasures object, with only the sentiment measures required, including updated information and statistics, but the original sentiment scores data.table untouched.

Author(s)

Samuel Borms

Examples

```
data("usnews")
data("lexicons")
data("valence")

# construct a sentomeasures object to start with
corpus <- sento_corpus(corpusdf = usnews)
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
ctr <- ctr_agg(howTime = c("equal_weight", "linear"), by = "year", lag = 3)
sentomeasures <- sento_measures(corpus, l, ctr)

# different selections
sel1 <- select_measures(sentomeasures, c("equal_weight"))
sel2 <- select_measures(sentomeasures, c("equal_weight", "linear"), do.combine = FALSE)
sel3 <- select_measures(sentomeasures, c("linear", "LM_eng"))
sel4 <- select_measures(sentomeasures, c("linear", "LM_eng", "wsj", "economy"),
                        do.combine = FALSE)
sel5 <- select_measures(sentomeasures, c("linear", "LM_eng"),
                        dates = "date >= '1989-12-31' & date <= '2000-12-31'")
d <- seq(as.Date("1980-01-01"), as.Date("2013-12-01"), by = "month")
sel6 <- select_measures(sentomeasures, c("linear", "LM_eng"), dates = d)
```

sentometrics

An Integrated Framework for Textual Sentiment Time Series Aggregation and Prediction

Description

The sentometrics package is designed to do time series analysis based on textual sentiment. It accounts for the intrinsic challenge that, for a given text, sentiment can be computed in hundreds of different ways, as well as the large number of possibilities to pool sentiment across text and time. This additional layer of manipulation does not exist in standard time series analysis packages. As a final outcome, this package provides an automated means to econometrically model the impact of sentiment in texts on a given variable, by first computing a wide range of textual sentiment time series and then selecting the sentiment times series that are most informative. The package created therefore integrates the qualification of sentiment from texts, the aggregation into different sentiment measures and the optimized prediction based on these measures.

Main functions

- Sentiment computation and aggregation into sentiment measures: [sento_corpus](#), [ctr_agg](#), [compute_sentiment](#), [sento_measures](#), [to_global](#)
- Sparse modelling: [ctr_model](#), [sento_model](#)
- Prediction and post-modelling analysis: [predict.sentomodel](#), [retrieve_attributions](#), [perform_MCS](#)

Update

The latest version of the package is available at <https://github.com/ArdiaD/Sentometrics>.

Note

The ideas behind the sentiment aggregation framework can be consulted in the working paper titled “Questioning the news about economic growth: Sparse forecasting using thousands of news-based sentiment values” (Ardia, Bluteau & Boudt, 2017) at https://papers.ssrn.com/sol3/papers.cfm?abstract_id=2976084.

Please cite the package in publications. Use `citation("sentometrics")`.

Author(s)

Samuel Borms, Keven Bluteau, David Ardia and Kris Boudt.

sento_corpus

Create a sentocorpus object

Description

Formalizes a collection of texts into a well-defined corpus object, by calling the `quanteda::corpus()` function. The **quanteda** package is a (very) fast text mining package; for more info, see [quanteda](#). This function mainly performs a set of checks on the input data and prepares the corpus for further sentiment analysis.

Usage

```
sento_corpus(corpusdf, do.clean = FALSE)
```

Arguments

<code>corpusdf</code>	a <code>data.frame</code> with as named columns and <i>in this order</i> : a document id column, a date column, a text column (i.e. the columns where all texts to analyze reside), and a series of feature columns of type <code>numeric</code> , with values pointing to the applicability of a particular feature to a particular text. The latter columns are often binary (1 means the feature is applicable to the document in the same row) or as a percentage to specify the degree of connectedness of a feature to a document. Features could be topics (e.g. legal, political or economic), but also article sources (e.g. online or printed press), amongst many more options. If no particular features are of interest to your analysis, have only one additional column with all values set to 1. Provide the date column as "yyyy-mm-dd". The id column should be in character mode. All spaces in the names of the features are automatically replaced by underscores.
<code>do.clean</code>	a logical, if TRUE all texts undergo a cleaning routine to eliminate common textual garbage. This includes a brute force replacement of HTML tags and non-alphanumeric characters by an empty string.

Details

A sentocorpus object can be regarded as a specialized instance of a **quanteda** corpus. In theory, all **quanteda** functions applicable to its corpus object can also be applied to a sentocorpus object. However, changing a given sentocorpus object too drastically using some of **quanteda**'s functions might alter the very structure the corpus is meant to have (as defined in the "corpusdf" argument) to be able to be used as an input in other functions of the **sentometrics** package. There are functions, including `quanteda::corpus_sample()` or `quanteda::corpus_subset()`, that do not change the actual corpus structure and may come in handy. To add additional features, we recommend to use `add_features()`. In the future, we will formalize the interaction between the **quanteda** package (as well as other text mining packages).

Value

A sentocorpus object, derived from a **quanteda** corpus classed list keeping the elements documents, metadata and settings. The documents element incorporates the corpus represented as a `data.frame`.

Author(s)

Samuel Borms

See Also

[corpus](#)

Examples

```
data("usnews")

# corpus construction
corpus <- sento_corpus(corpusdf = usnews)

# take a random subset using a quanteda's package function
corpusSmall <- quanteda::corpus_sample(corpus, size = 500)
```

sento_measures

One-way road towards a sentomeasures object

Description

Wrapper function which assembles calls to `compute_sentiment()` and `perform_agg()`, and includes the input sentocorpus and computed sentiment scores in its output. Serves as the most direct way towards a panel of textual sentiment measures, and a sentomeasures object.

Usage

```
sento_measures(sentocorpus, lexicons, ctr)
```

Arguments

<code>sentocorpus</code>	a sentocorpus object.
<code>lexicons</code>	output from a <code>setup_lexicons()</code> call.
<code>ctr</code>	output from a <code>ctr_agg()</code> call.

Value

A sentomeasures object, which is a list containing:

measures	a data.table with a date column and all textual sentiment measures as remaining columns.
features	a character vector of the different features.
lexicons	a character vector of the different lexicons used.
time	a character vector of the different time weighting schemes used.
by	a single character vector specifying the time interval of aggregation used.
stats	a data.frame with a series of elementary statistics (mean, standard deviation, maximum, minimum, and average correlation with all other measures) for each individual sentiment measure.
sentiment	the sentiment scores data.table with a date and lexicon–feature sentiment scores columns.
howWithin	a character vector to remind how sentiment within documents was aggregated.
howDocs	a character vector to remind how sentiment across documents was aggregated.
fill	a single character vector that specifies if and how missing dates have been added before aggregation across time was carried out.
attribWeights	a list of document weights and time weights that are extracted in a call to retrieve_attributions(). Serves further no direct purpose.

Author(s)

Samuel Borms, Keven Bluteau

See Also

[compute_sentiment](#), [perform_agg](#)

Examples

```
data("usnews")
data("lexicons")
data("valence")

# construct a sentomeasures object to start with
corpus <- sento_corpus(corpusdf = usnews)
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
ctr <- ctr_agg(howWithin = "tf-idf",
               howDocs = "proportional",
               howTime = c("equal_weight", "linear", "almon"),
               by = "month",
               lag = 3,
               ordersAlm = 1:3,
               do.inverseAlm = TRUE,
               do.normalizeAlm = TRUE)
sento_measures <- sento_measures(corpus, l, ctr)
summary(sento_measures)
```

sento_model

*Optimized and automated sparse regression***Description**

Linear or nonlinear penalized regression of any dependent variable on the wide number of sentiment measures and potentially other explanatory variables. Either performs a regression given the provided variables at once, or computes regressions sequentially for a given sample size over a longer time horizon, with associated one-step ahead prediction performance metrics.

Usage

```
sento_model(sentomeasures, y, x = NULL, ctr)
```

Arguments

sentomeasures	a sentomeasures object. There should be at least two explanatory variables including the ones provided through the x argument.
y	a one-column data.frame or a numeric vector capturing the dependent (response) variable. In case of a logistic regression, the response variable is either a factor or a matrix with the factors represented by the columns as binary indicators, with the second factor level or column as the reference class in case of a binomial logistic regression. No NA values are allowed.
x	a named data.frame with other explanatory variables as numeric, by default set to NULL.
ctr	output from a ctr_model() call.

Details

Models are computed using the elastic net regularization as implemented in the **glmnet** package, to account for the multidimensionality of the sentiment measures. Additional explanatory variables are not subject to shrinkage. Independent variables are normalized in the regression process, but coefficients are returned in their original space. For a helpful introduction to **glmnet**, we refer to the [vignette](#). The optimal elastic net parameters lambda and alpha are calibrated either through a to specify information criterion or through cross-validation (based on the "rolling forecasting origin" principle). In the latter case, the training metric is automatically set to "RMSE" for a linear model and to "Accuracy" for a logistic model. We suppressed many of the details that can be supplied to the `glmnet::glmnet()` and `caret::train()` functions we rely on for estimation and calibration through cross-validation, for the sake of user-friendliness.

Value

If `ctr$do.iter == FALSE`, a sentomodel object which is a list containing:

reg	optimized regression, i.e. a model-specific glmnet object.
model	the input argument ctr\$model, to remind of the type of model that was estimated.
x	the matrix of the values used in the regression for all explanatory variables.
alpha	optimized calibrated alpha.
lambda	optimized calibrated lambda.

trained	output from <code>caret::train</code> call (if <code>ctr\$type == "cv"</code>).
ic	a list composed of two elements: the information criterion used in the calibration under "criterion", and a vector of all minimum information criterion values for each value in <code>alphas</code> under "opts" (if <code>ctr\$type != "cv"</code>).
dates	sample reference dates, being the earliest and most recent date from the <code>sentomeasures</code> object accounted for in the estimation window.
nVar	the sum of the number of sentiment measures and other explanatory variables inputted.
discarded	a named logical vector of length equal to the number of sentiment measures, in which TRUE indicates that the particular sentiment measure has not been considered in the regression process.

If `ctr$do.iter == TRUE`, a `sentomodeliter` object which is a list containing:

models	all sparse regressions, i.e. separate <code>sento_model</code> objects as above, as a list with as names the dates from the perspective of the sentiment measures at which predictions for performance measurement are carried out (i.e. one date step beyond the date found at <code>sento_model\$date</code>).
alphas	optimized calibrated alphas.
lambdas	optimized calibrated lambdas.
performance	a <code>data.frame</code> with performance-related measures, being "RMSFE" (root mean squared forecasting error), "MAD" (mean absolute deviation), "MDA" (mean directional accuracy, in which's calculation zero is considered as a positive; in percentage points), "accuracy" (proportion of correctly predicted classes in case of a logistic regression; in percentage points), and each's respective individual values in the sample. Directional accuracy is measured by comparing the change in the realized response with the change in the prediction between two consecutive time points (omitting the very first prediction, resulting in NA). Only the relevant performance statistics are given depending on the type of regression. Dates are as in the "models" output element, i.e. from the perspective of the sentiment measures.

Author(s)

Samuel Borms, Keven Bluteau

See Also

[ctr_model](#), [glmnet](#), [train](#), [retrieve_attributions](#)

Examples

```
data("usnews")
data("lexicons")
data("valence")
data("epu")

# construct a sentomeasures object to start with
corpusAll <- sento_corpus(corpusdf = usnews)
corpus <- quanteda::corpus_subset(corpusAll, date >= "1980-01-01" & date < "2014-10-01")
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
ctr <- ctr_agg(howWithin = "tf-idf", howDocs = "proportional",
              howTime = c("equal_weight", "linear", "almon"),
```

```

        by = "month", lag = 3, ordersAlm = 1:3,
        do.inverseAlm = TRUE, do.normalizeAlm = TRUE)
sentomeasures <- sento_measures(corpus, 1, ctr)

# prepare y and other x variables
y <- epu[epu$date >= sentomeasures$measures$date[1], ]$index
length(y) == nrow(sentomeasures$measures) # TRUE
x <- data.frame(runif(length(y)), rnorm(length(y))) # two other (random) x variables
colnames(x) <- c("x1", "x2")

# a list with models based on the three implemented information criteria
out1 <- list()
for (ic in c("BIC", "AIC", "Cp")) {
  ctrIC <- ctr_model(model = "gaussian", type = ic, do.iter = FALSE, h = 0)
  out1[[ic]] <- sento_model(sentomeasures, y, x = x, ctr = ctrIC)
}

# some post-analysis (attribution)
attributions1 <- retrieve_attributions(out1[["Cp"]], sentomeasures)
attributions2 <- retrieve_attributions(out1[["Cp"]], sentomeasures,
                                     refDates = sentomeasures$measures$date[20:40])

# a cross-validation based model
ctrCV <- ctr_model(model = "gaussian", type = "cv", do.iter = FALSE,
                  h = 0, alphas = c(0.10, 0.50, 0.90), trainWindow = 350,
                  testWindow = 40, oos = 0, do.progress = TRUE)
out2 <- sento_model(sentomeasures, y, x = x, ctr = ctrCV)
summary(out2)

# a cross-validation based model but for a binomial target
yb <- epu[epu$date >= sentomeasures$measures$date[1], ]$above
ctrCVb <- ctr_model(model = "binomial", type = "cv", do.iter = FALSE,
                   h = 0, alphas = c(0.10, 0.50, 0.90), trainWindow = 350,
                   testWindow = 40, oos = 0, do.progress = TRUE)
out3 <- sento_model(sentomeasures, yb, x = x, ctr = ctrCVb)
summary(out3)

# an example of an iterative analysis
ctrIter <- ctr_model(model = "gaussian", type = "BIC", do.iter = TRUE,
                   h = 0, nSample = 300, start = 106)
out <- sento_model(sentomeasures, y, x = x, ctr = ctrIter)
summary(out)

# the same iterative analysis, parallelized
## Not run:
cl <- makeCluster(detectCores() - 2)
registerDoParallel(cl)
ctrIter <- ctr_model(model = "gaussian", type = "BIC", do.iter = TRUE,
                   h = 0, nSample = 300, start = 106, do.parallel = TRUE)
out <- sento_model(sentomeasures, y, x = x, ctr = ctrIter)
stopCluster(cl)
summary(out)
## End(Not run)

# some more post-analysis (attribution and prediction)
attributions <- retrieve_attributions(out, sentomeasures)
plot_attributions(attributions, "lexicons")

```

```

plot_attributions(attributions, "features")
plot_attributions(attributions, "time")

nx <- ncol(sentomeasures$measures) - 1 + ncol(x) # don't count date column
newx <- runif(nx) * cbind(sentomeasures$measures[, -1], x)[1:nrow(x), ]
preds2 <- predict(out2, newx = as.matrix(newx), type = "link")
preds3 <- predict(out3, newx = as.matrix(newx), type = "class")

```

setup_lexicons

Set up lexicons (and valence word list) for use in sentiment analysis

Description

Structures provided lexicons and potentially integrates valence words. One can also provide (part of) the built-in lexicons from `data("lexicons")` or a valence word list from `data("valence")` as an argument. Makes use of the `as_key()` function from the **sentimentr** package to make the output coherent and check for duplicates.

Usage

```
setup_lexicons(lexiconsIn, valenceIn = NULL, do.split = FALSE)
```

Arguments

lexiconsIn	a list of (raw) lexicons, each element being a <code>data.table</code> or <code>data.frame</code> with respectively a words column and a polarity score column. The lexicons should be appropriately named for clarity in terms of subsequently obtained sentiment measures. Alternatively, a subset of the already formatted built-in lexicons accessible via <code>lexicons</code> can be declared too, as part of the same list input. If only (some of) the package built-in lexicons want to be used, one can simply supply <code>lexicons[c(...)]</code> as an argument to either <code>sento_measures()</code> or <code>compute_sentiment()</code> . However, it is strongly recommended to pass the lexicons (and a valence word list) that want to be used through this function.
valenceIn	a single valence word list as a <code>data.table</code> or <code>data.frame</code> with respectively a words column, a type column (1 for negators, 2 for amplifiers/intensifiers, and 3 for deamplifiers/downtoners) and a score column. Suggested scores are -1, 2 and 0.5 respectively, and should be the same within each type. Alternatively, this argument can be one of the already formatted built-in valence word lists accessible via <code>valence</code> . If <code>NULL</code> , no valence word list is part of the output.
do.split	a logical that if <code>TRUE</code> splits every lexicon into a separate positive polarity and negative polarity lexicon.

Value

A list with each lexicon as a separate element according to its name, as a `data.table`, and a list element named `valence` that comprises the valence words. Every `x` column contains the words, every `y` column contains the polarity score, and for the valence word list, `t` contains the word type. If a valence word list is provided, all lexicons are expanded by copying the respective lexicon, and changing the words and scores according to the valence word type: "NOT_" is added for negators, "VERY_" is added for amplifiers and "HARDLY_" is added for deamplifiers. Lexicon scores are multiplied by -1, 2 and 0.5 by default, respectively, or the first value of the scores column of the valence word list.

Author(s)

Samuel Borms

See Also

[as_key](#)

Examples

```
data("lexicons")
data("valence")

# sets up output list straight from built-in word lists including valence words
l1 <- c(lexicons[c("LM_eng", "HENRY_eng")], valence[["eng"]])

# using function, including a self-made lexicon, with and without valence shifters
lexIn <- c(list(myLexicon = data.table(w = c("nice", "boring"), s = c(2, -1))),
           lexicons[c("GI_eng")])
valIn <- valence[["valence_eng"]]
l2 <- setup_lexicons(lexIn)
l3 <- setup_lexicons(lexIn, valIn)
l4 <- setup_lexicons(lexIn, valIn, do.split = TRUE)
```

sp500

Monthly S&P 500 Index returns

Description

Monthly returns for the S&P 500 Index between March 1988 and December 2014, including a binomial and a multinomial example series. It has following columns:

- date. Date as "yyyy-mm-01".
- return. A numeric value as the return that was achieved during the corresponding month; for example the first return value is the price change from beginning of March 1988 to beginning of April 1988.
- up. A factor with value "pos" if the return is greater than zero, else "neg".
- upMulti. A factor with values "pos+", "pos", "neg" and "neg-" if returns are greater than 0.05 and 0, or smaller than 0 and -0.05, respectively and in a mutually exclusive sense.

Usage

```
data("sp500")
```

Format

A data.frame with 322 rows and 4 columns.

Source

[S&P 500 \(^GSPC\) at Yahoo Finance](#)

`to_global`*Merge sentiment measures into one global sentiment measure*

Description

Merges all sentiment measures into one global textual sentiment measure based on a set of weights to indicate the importance of each component in the lexicons, features and time vectors as specified in the input `sentomeasures` object. The global measure is composed as the multiplication of the individual weights across the three dimensions times the sentiment value per date observation.

Usage

```
to_global(sentomeasures, lex = 1, feat = 1, time = 1)
```

Arguments

<code>sentomeasures</code>	a <code>sentomeasures</code> object.
<code>lex</code>	a numeric vector of weights, of size <code>length(sentomeasures\$lexicons)</code> , in the same order and summing to one. By default set to 1, which means equally weighted.
<code>feat</code>	a numeric vector of weights, of size <code>length(sentomeasures\$features)</code> , in the same order and summing to one. By default set to 1, which means equally weighted.
<code>time</code>	a numeric vector of weights, of size <code>length(sentomeasures\$time)</code> , in the same order and summing to one. By default set to 1, which means equally weighted.

Details

This function returns no `sentomeasures` object, however the global sentiment measure as outputted can be added to regressions as an additional variable using the `x` argument in the `sento_model()` function.

Value

A `data.frame` with the values for the global sentiment measure under the `global` column and dates as row names.

Author(s)

Samuel Borms

See Also

[sento_model](#)

Examples

```
data("usnews")
data("lexicons")
data("valence")

# construct a sentomeasures object to start with
corpus <- sento_corpus(corpusdf = usnews)
l <- setup_lexicons(lexicons[c("LM_eng", "HENRY_eng")], valence[["valence_eng"]])
ctr <- ctr_agg(howTime = c("equal_weight", "linear"), by = "year", lag = 3)
sentomeasures <- sento_measures(corpus, l, ctr)

# merge into one global sentiment measure, with specified weighting for lexicons and features
global <- to_global(sentomeasures, lex = c(0.40, 0.60),
                    feat = c(0.10, 0.20, 0.30, 0.40),
                    time = 1)
```

usnews	<i>Texts relevant (and not) to the US economy</i>
--------	---

Description

A collection of texts annotated by humans in terms of relevance to the US economy or not. The texts come from two major journals in the US (The Wall Street Journal and The Washington Post) and cover 6801 documents between 1980 and 2014. It contains following information:

- id. a character ID identifier.
- date. Date as "yyyy-mm-dd".
- text. Texts in character format.
- wsj. Equals 1 if the article comes from The Wall Street Journal.
- wapo. Equals 1 if the article comes from The Washington Post.
- economy. Equals 1 if the article is relevant to the US economy.
- noneconomy. Equals 1 if the article is not relevant to the US economy.

Usage

```
data("usnews")
```

Format

A `data.frame`, formatted as required to be an input for `sento_corpus`.

Source

Economic News Article Tone and Relevance

Examples

```
data("usnews")
usnews[3192, "text"]
usnews[1:5, c("id", "date", "text")]
```

valence	<i>Built-in valence word lists</i>
---------	------------------------------------

Description

A list containing all built-in valence word lists, a `data.table` with three columns: a `x` column with the words, a `t` column with the type of valence words, and a `y` column with the value associated to a word and type of valence shifter. The list element names incorporate the language of the valence word list. All non-English word lists are translated. The valence word lists are in the form required for further sentiment analysis. The built-in valence word lists are the following:

- `valence_eng`
- `valence_fr`
- `valence_nl`

Usage

```
data("valence")
```

Format

A list with all built-in valence word lists, appropriately named.

Source

[hash_valence_shifters](#) (negators)

Index

*Topic **datasets**

- epu, [10](#)
- lexicons, [13](#)
- sp500, [31](#)
- usnews, [33](#)
- valence, [34](#)

add_features, [2](#)

almons, [3](#)

as_key, [31](#)

compute_sentiment, [4](#), [6](#), [15](#), [23](#), [26](#)

corpus, [25](#)

ctr_agg, [3](#), [5](#), [11](#), [12](#), [15](#), [23](#)

ctr_merge, [7](#), [14](#)

ctr_model, [8](#), [23](#), [28](#)

dfm, [5](#)

epu, [10](#)

exponentials, [11](#)

fill_measures, [6](#), [11](#)

get_hows, [12](#)

glmnet, [28](#)

hash_valence_shifters, [34](#)

lexicons, [13](#)

MCSprocedure, [16](#)

merge_measures, [8](#), [14](#)

perform_agg, [15](#), [26](#)

perform_MCS, [16](#), [23](#)

plot.sentomeasures, [16](#)

plot.sentomodeliter, [17](#)

plot_attributions, [19](#)

predict.glmnet, [20](#)

predict.sentomodel, [19](#), [23](#)

retrieve_attributions, [20](#), [23](#), [28](#)

scale, [21](#)

scale.sentomeasures, [21](#)

select_measures, [22](#)

sento_corpus, [23](#), [24](#), [33](#)

sento_measures, [15](#), [23](#), [25](#)

sento_model, [8–10](#), [16](#), [19–21](#), [23](#), [27](#), [32](#)

sentometrics, [23](#)

sentometrics-package (sentometrics), [23](#)

setup_lexicons, [30](#)

sp500, [31](#)

to_global, [23](#), [32](#)

tokenize, [5](#)

train, [28](#)

usnews, [33](#)

valence, [34](#)