

The **sentometrics** R package to compute, aggregate and predict with textual sentiment

David Ardia
University of Neuchâtel
Laval University

Keven Bluteau
University of Neuchâtel
Vrije Universiteit Brussel

Samuel Borms
University of Neuchâtel
Vrije Universiteit Brussel

Kris Boudt
Vrije Universiteit Brussel
Vrije Universiteit Amsterdam

Abstract

This vignette is a hands-on introduction to optimized textual sentiment indexation using the **sentometrics** package. The package allows to associate sentiment to many texts at once, aggregate the textual sentiment scores into multiple time series, and use these time series to predict any other variable. Sentiment analysis of texts is increasingly used to capture the information value of qualitative data in relation to quantitative variables. The workflow of the package is illustrated using a built-in corpus of texts from two major U.S. journals to predict an index of economic policy uncertainty.

Keywords: sentometrics, textual sentiment, R, time series, penalized regression.

1. Introduction

The sentiment transmitted through texts undoubtedly is a prime candidate driver of many variables, given the immense flow at which texts are spread out and analysed. Texts, qualitative by nature, need to be quantified first before they can be related to other quantitative data. The most popular means to do so is by assigning a sentiment score that represents the polarity of a text. This is called *textual sentiment*. Over the past decade, the number of academic articles that transform texts into quantitative variables using this approach has steadily grown. The creation of lexicon-based news measures to forecast stock market volatility (Caporin and Poli 2017), the study of the tone in earnings press releases (Davis *et al.* 2012) and the analysis of textual sentiment to predict stock returns (Tetlock *et al.* 2008) are convincing examples of the added value of computationally studying texts in economics and finance. The application of textual (sentiment) analysis vastly exceeds these two fields however, equally of interest to domains as diverse as marketing, politics and history.

The **sentometrics** software package positions itself as both supplementary and integrative with respect to the current text mining R universe. It is limited in terms of pure text mining capabilities, for which **cleanNLP** (Arnold 2017), **quanteda** (Benoit *et al.* 2017), **tidytext** (Silge and Robinson 2016), **qdap** (Rinker 2013) and **tm** (Feinerer *et al.* 2008) are more than equipped already. The acronym YATAP¹ does not apply here. On the other hand, the package is not a fully fledged statistical toolbox either. Better, it combines the strengths of text mining and econometric analysis, bringing forward a well-defined modelling workflow, specifically applied to studying the evolution of textual sentiment and its impact on other quantities. The workflow as defined was brought forward by Ardia *et al.* (2017), who show that the incorporation of sentiment in texts allows to better predict Germany's industrial production growth. It extends most, if not all, contemporary text mining workflows by (i) stretching out fine-grained textual sentiment into the construction of various sentiment time series (going by the synonyms of measures or indices), and (ii) using these measures in a flexible prediction setup. The **sentometrics** package greatly benefits from three major packages: **quanteda** for corpus construction and corpus manipulation for sentiment analysis, **data.table** (Dowle and Srinivasan 2017) for the aggregation of textual sentiment into time series, and especially **glmnet** (Friedman *et al.* 2010) but also **caret** (Kuhn 2017) for sparse model estimation.

2. The framework behind sentometrics

Following Ardia *et al.* (2017) but grouping together the different phases according to the package's structure, there are five main steps in the workflow. The steps are explained below, together with the functions from the **sentometrics** package that relate to the specific step. How to use the functions is demonstrated in the next section. We minimize the mathematical details to clarify the exposition and stay close to the actual implementation in the package.

Step 1: Acquire and pre-process a selection of texts and generate relevant features

The first step consists in the acquisition of a corpus of texts, of any size and from any source available. Texts could be crawled from the web or retrieved from news databases. A minimal requirement is that every text has a timestamp associated. These texts are ideally cleaned

¹Yet Another Text Analysis Package.

such that graphical and web-related elements (for example, HTML tags) are removed. This results in a set of documents $d_{n,t}$ for $n = 1, \dots, N_t$ and time points $t = 1, \dots, T$, where N_t is the total number of documents at time t .

Secondly, *features* have to be defined and mapped to the documents. Features can come in many forms: news sources, entities (individuals, companies or countries discussed) or text topics. The mapping implicitly permits to subdivide the corpus into many smaller groups with a common interest. Many data providers tend to enrich their textual data with information that can be used as features. If this is not the case, topic modelling or entity recognition techniques are valid alternatives. Human classification or keyword(s) occurrence searches are other, more simple, options. The extraction and inclusion of features is an important part of the analysis and should be related to the variable that is meant to be predicted. The texts and features have to be structured in a rectangular fashion. Every row is a document that is mapped to the features through numerical values $w_{n,t,k}$. The features are denoted by $k = 1, \dots, K$. The values are indicative of the relevance of a feature to a document, and are often binary.

When the rectangular data structure is set up, it needs to be passed to the `sentocorpus()` function to create a more formal corpus, resulting in a *sentocorpus* object. The reason for this intermediate step is twofold: it controls whether all corpus requirements for further analysis are met, and it allows to perform more efficient operations on the corpus both for corpus manipulation and sentiment calculation. The `add_features()` function can be used to add or generate new features, as shown in the illustration. When the corpus is constructed, it is up to the user to decide which texts have to be kept for the actual sentiment analysis.

Step 2: Choose lexicons and compute textual sentiment

For every text, sentiment is to be computed. There are many ways to do so, from straightforward to very complex. The approach taken here is to augment the *bag-of-words* model by including *valence shifters*. The bag-of-words model looks for words (*unigrams*) in a text that are included in a pre-defined word list of polarized (positive and negative) words. Textual sentiment is then the summation of all scores associated to these detected words considering their frequency and given a within-document weighting scheme. Incorporating valence shifters goes one step beyond, taking a *bigram* perspective. The impact of the word appearing before the detected word is evaluated as well.² A common example is ‘not good’, which under the default approach would get a score of 1 (‘good’), but now ends up having a score of -1 due to the negator ‘not’.

Before being able to apply the enhanced bag-of-words model, a series of lexicons and one list of valence shifters need to be determined. The *sentometrics* package benefits from built-in word lists in English, French and Dutch, with the latter two often as a translation.³ Defining the lexicons is done through the `setup_lexicons()` function.

For every lexicon $l = 1, \dots, L$, each document $d_{n,t}$ gets assigned a sentiment score $s_{n,t,l}$. These scores are subsequently multiplied by the feature weights to obtain lexicon- and feature-

²Valence shifters can also appear in positions other than right before a certain word. The *sentimentr* package (Rinker 2017) does a good job at accounting for this. We decided not to implement this layer of complexity, as it would take too long to compute sentiment for tens of thousands of documents at once. It is kept for a future version of the package.

³However, all lexicons are applied to the corpus, without discrimination based on language.

specific sentiment scores, as $s_{n,t,l,k} = s_{n,t,l} \times w_{n,t,k}$, with k the index denoting the feature. If the document does not correspond to the feature, the value of $s_{n,t,l,k}$ is zero.

The sentiment calculation is performed with `compute_sentiment()`. This function outputs a `data.table` with all values for $s_{n,t,l,k}$, which serves as the basis for aggregation in the next step.

Step 3: Aggregate into textual sentiment time series

The purpose in this step is to aggregate the individual sentiment scores and obtain various representative time series. Two main aggregations are performed. The first collapses all sentiment scores across documents within the same frequency (*e.g.*, day or month, as defined by t) into one score. The weighted sum that does so is $s_{t,l,k} = \sum_{n=1}^{N_t} \theta_n s_{n,t,l,k}$, where the weights θ_n define the importance of each document n at time t (for instance, based on the length of the text). The second smooths the newly aggregated sentiment scores over time, as $s_{u,l,k,b} = \sum_{t=t_\tau}^u b_t s_{t,l,k}$, where $t_\tau = u - \tau + 1$. The time weighting schemes $b = 1, \dots, B$ go with different values for b_t to smooth the time series, with a time lag of τ . The first $\tau - 1$ observations are dropped from the original time indices, such that $u = \tau, \dots, T$ becomes the time index for the ultimate time series. This leaves $N = T - \tau + 1$ time series observations.

The entire aggregation setup is specified by means of the `ctr_agg()` function, including the within-document aggregation needed for the sentiment analysis. The `sentomeasures()` function performs both the sentiment calculation (via `compute_sentiment()`) and time series aggregation (via `perform_agg()`), outputting a *sentomeasures* object. The number of obtained time series is equal to $P = L \times K \times B$. Every time series covers one aspect of each dimension; an example is the textual sentiment as computed by lexicon l for feature k aggregated across time using scheme b . A single time series incorporating such combination is designated by $s_{u,p}$ across all values of u , for $p = 1, \dots, P$. The ensemble of time series captures both different information and the same information in different ways.

Step 4: Calibrate sparse regression model and perform (out-of-sample) predictions

The sentiment measures are now regular time series variables that can be applied in regressions. In case of a linear regression, the reference equation is:

$$y_{u+h} = \delta + \gamma' x_u + \beta_1 s_{u,1} + \dots + \beta_p s_{u,p} + \dots + \beta_P s_{u,P} + \epsilon_{u+h}.$$

The target variable y_{u+h} is often a variable to forecast, *i.e.*, $h > 0$. Logistic regression is available in the package based on the same underlying linear structure. Other (non-sentiment) variables are captured by x_u and γ is the associated parameter vector. All other variables are the textual sentiment variables as constructed before. Let $\beta = (\beta_1, \dots, \beta_P)'$.

The large dimensionality of P in itself and relative to the number of observations poses a problem to regular least squares regression. Instead, estimation through a penalized regression relying on the elastic net regularization of [Zou and Hastie \(2005\)](#) is more appropriate. Regularization, in short, shrinks the coefficients of the least informative variables towards zero. The estimates of the model coefficients for the textual sentiment indices are in $\hat{\beta}$, usually a sparse vector, depending on the severity of the shrinkage. Elastic net solves an optimization problem that requires two parameters as an input. The parameter `alpha` defines the trade-off between the Ridge ([Hoerl and Kennard 1970](#)) and the LASSO ([Tibshirani 1996](#))

regularization, respectively when it is equal to 0 and 1. The `lambda` parameter defines the level of regularization. The two parameters are calibrated such that they are optimal to the regression equation at hand. The **sentometrics** package allows four ways to do the calibration, either through cross-validation, or based on three information criteria adapted to elastic net (Tibshirani and Taylor 2012).

An analysis of typical interest is the sequential estimation of a regression model and out-of-sample prediction. For a given sample size $m < N$, a regression model is estimated for the first m observations and used to predict the target variable one-step ahead, *i.e.*, at $m + 1$. This procedure is repeated which gives a series of estimates, to be compared with the realized values to assess the average out-of-sample prediction performance.

The type of model, the calibration approach and other modelling decisions are defined via the `ctr_model()` function. The (iterative) model estimation and calibration is performed by calling `sento_model()`. The output is a *sentomodel* (one model) or a *sentomodeliter* (a collection of iteratively estimated *sentomodel* objects and out-of-sample predictions) object.

Step 5: Evaluate prediction performance and retrieve sentiment attributions

The *sentomodeliter* object carries an overview of performance measures relevant to the type of model estimated. Additionally, it can easily be plotted to visually inspect the prediction precision. A good way to compare different models, sentiment-based or not, is to construct a model confidence set (Hansen *et al.* 2011). This so-called set isolates the models that are statistically the best in terms of predictive ability, within a confidence level. A wrapper around the core function of the **MCS** package (Catania and Bernardi 2017) is made available in the **sentometrics** package, `perform_MCS()`, to carry out this comparative analysis.

The aggregation into textual sentiment time series is entirely linear. This allows, based on the estimated coefficients $\hat{\beta}$, to compute every corresponding underlying dimension's sentiment *attribution* to a given prediction. For example, the attribution of a certain feature k in the forecast of the target variable, at a particular date, is the weighted sum of the model coefficients and the values of the sentiment measures in which k plays a role. In a similar vein, the attribution can be computed for all features, lexicons, time weighting schemes and individual documents. Through attribution, a prediction is broken down in its respective components. The functions `retrieve_attributions()` and `plot_attributions()` provide insights into the evolution of the prediction attribution along the different dimensions.

3. The **sentometrics** R package

The **sentometrics** package takes care of *Steps 2–5* in full. The texts collection in *Step 1* is, obviously, left to the user. The other parts (cleaning, feature generation and selection) are also mostly left to the user, but support for these tasks is present within the package. In what follows, an elaborate example shows how to put the steps into practice using the **sentometrics** package. We first explain the two core aggregation functions, `ctr_agg()` and `ctr_model()`. Then we illustrate the most straightforward workflow using built-in data. To end, we briefly touch upon other handy functions that are present in the package.

3.1. Deep-dive into the two main control functions

Aggregation control

The required arguments of the `ctr_agg()` function are the following: `howWithin`, `howDocs`, `howTime`, `do.ignoreZeros`, `by`, `lag` and `fill`. The optional arguments are `alphasExp`, `ordersAlm`, `do.inverseAlm`, `do.normalizeAlm`, `weights` and `dfm`. We focus the explanation on the three first arguments, central to the aggregation:

- **howWithin** – This argument defines how sentiment is aggregated within the same document. Do you simply want to take the difference between the number of positive and negative words? Do you prefer to do the same but normalize the sentiment score based on the total number of words within the document? Or do you think that words appearing very frequently across all documents should be down-weighted? Depending on what you prefer, you can respectively select the "counts", "proportional" or "tf-idf" option. The last is an abbreviation for term frequency-inverse document frequency, a well-known concept within computational linguistics.
- **howDocs** – This argument defines how sentiment is aggregated across all documents at the same date (or frequency), that is, it sets the weights θ_n . There are two options: either "equal_weight", which gives the same weight to every document, or "proportional", which gives higher weights to documents with more words, relative to the document population at a given date. The `do.ignoreZeros` argument can be used to ignore documents with zero sentiment in the computation of the across-document weights. By default these documents are overlooked. This primarily avoids the incorporation of documents not relevant to a particular feature (as in those cases $s_{n,t,l,k}$ is exactly zero, because $w_{n,t,k} = 0$), which could lead to a bias of sentiment towards zero.
- **howTime** – This argument defines how sentiment is aggregated across dates, to smoothen the sentiment time series and acknowledge that sentiment at a given point is at least partly based on sentiment and information from the past. The `lag` argument has the role of τ ; it dictates how far to go back. The "equal_weight" option is similar to a simple weighted moving average, "linear" and "exponential" are two options which give weights to the observations according to a linear or an exponential curve, and "almon" does the same but based on so-called Almon polynomials. The last two curves have respective arguments to define their shape(s), being `alphasExp`, and `ordersAlm`, `do.inverseAlm` and `do.normalizeAlm`. If desired, user-constructed weights can be supplied via `weights` as a named `data.frame`. All the weighting schemes define the different b_t values. The `fill` argument is of sizeable importance here. It is used to add in dates for which not a single document was available. These added, originally missing, dates are given a value of 0 ("zero") or the most recent value ("latest"). The option "none" will not fill up the date sequence at all. Adding in dates (or not) impacts the time aggregation by respectively combining the latest *consecutive* dates, or the latest *available* dates.

Via the `by` argument, the time interval at which the time series have to be aggregated is chosen: daily ("day"), weekly ("week"), monthly ("month") or yearly ("year"). The `dfm` argument, short for document-feature matrix, is there to provide additional flexibility in

terms of which words are part of the universe the lexicons will look through to find polarized words. It can be safely ignored; more information is found in the documentation manual.

Modelling control

The `ctr_model()` function requires at least following arguments: `model`, `type`, `intercept`, `h`, `do.iter`, `alphas` and `do.progress`. The `model` argument can be "gaussian" (linear), "binomial" or "multinomial". The `type` specifies the calibration procedure to find the most appropriate `alpha` and `lambda`. It is either "cv" (cross-validation) or one of three information criteria ("BIC", "AIC" or "Cp"). The argument `alphas` can be altered to change the possible values for `alpha`. The possible values for `lambda` are generated internally (by the `glmnet::glmnet()` function call). To enact an iterative model estimation and a one-step ahead out-of-sample analysis, set `do.iter = TRUE`. The argument `h`, positive or negative, shifts the response variable according to y_{u+h} .⁴ The arguments `intercept` and `do.progress` respectively add an intercept to the model and print progress statements, or not.

The arguments `trainWindow`, `testWindow` and `oos` are needed when model calibration is performed through cross-validation, that is, when `type = "cv"`. As an example, assume there are 120 observations in total, `trainWindow = 80`, `testWindow = 10` and `oos = 5`. In the first round of cross-validation, a model is estimated for a certain `alpha` and `lambda` combination with the first 80 observations, then 5 observations are skipped, and predictions are generated for observations 86 to 95. The next round does the same but with all observations moved one step forward. This is done until the end of the total sample is reached and repeated for all possible parameter combinations, relying on the `caret::train()` function. The optimal `alpha` and `lambda` couple is the one that induced the lowest average prediction error. The cross-validation implemented is thus based on the "rolling forecasting origin" principle.

The arguments `nSample`, `start` and `do.parallel` are used for iterative modelling, when `do.iter = TRUE`. The first argument is the size of the sample that is used each time to re-estimate the model, *i.e.*, m . The second argument can be used to only run a (later) subset of the iterations; by default it is set to 1 which runs all iterations. The total number of iterations is equal to `length(y) - nSample - abs(h) - oos`, with `y` the response variable as a vector. If `do.parallel = TRUE` and a parallel backend is registered, the `%dopar%` construct from the **foreach** (Analytics and Weston 2015b) package is utilized to speed up the multiple model estimations altogether.

3.2. Illustration of the main workflow

As highlighted, the input texts have to be structured rectangularly, with every row a document. The `data.frame` structure fits this need. We demonstrate the workflow using the `esu` built-in dataset, a collection of news articles from The Wall Street Journal and The Washington Post between 1995 and 2014.⁵ The data is loaded below.

```
R> data("usnews")
R> class(usnews)
```

⁴Be careful, if the input response variable is not aligned date-wise with the sentiment measures and the other explanatory variables, `h` cannot be interpreted as the exact prediction horizon. In other words, `h` only shifts the input variables as they are provided, no more.

⁵The data originates from <https://www.crowdfunder.com/data-for-everyone> (under "Economic News Article Tone and Relevance").


```
[1] "data.frame"
```

The "id", "date" and "text" columns are mandatory. All other columns are reserved for features, of type `numeric`. There are four original features. The first two indicate the news source, the latter two the relevance of every document to the U.S. economy. The feature values $w_{n,t,k}$ are binary.

```
R> colnames(usnews)
```

```
[1] "id" "date" "text" "wsj" "wapo" "economy" "noneconomy"
```

To access the texts, simply do `usnews[["text"]]`. To put the texts and features into a corpus structure, call the `sento_corpus()` function. If you have no features available, the corpus can still be created without any feature columns in the input `data.frame`. In that situation, a dummy feature called "dummy" with a score of 1 for all texts is added to the *sentocorpus* output object.

```
R> corpusAll <- sento_corpus(usnews)
R> class(corpusAll)
```

```
[1] "sentocorpus" "corpus" "list"
```

The `sento_corpus()` function creates a *sentocorpus* object on top of the **quanteda**'s package *corpus* object. This means that many functions from **quanteda** to manipulate corpora can be applied to a *sentocorpus* object as well. This is helpful as can be seen below, where we limit the corpus to all articles between January 1995 and September 2014.

```
R> corpus <- quanteda::corpus_subset(corpusAll, date < "2014-10-1")
R> corpus
```

Corpus consisting of 4,097 documents and 5 docvars.

The docvars are the "date" column and all feature columns. To round off *Step 1*, we add two features related to uncertainty using `add_features()`. The features `el` and `war` give a score of 1 to documents in which respectively the words "election" and "war" appear, which turns out to be the case for 139 and 112 documents. This gives $K = 6$ features. The `add_features()` function comes in especially handy when the only feature present is the automatically created "dummy" feature.⁶

```
R> corpus <- add_features(corpus, keywords = list(el = "election", war = "war"))
R> c(sum(corpus$documents$el), sum(corpus$documents$war))
```

```
[1] 256 1486
```

⁶To delete a feature, one can do `quanteda::docvars(corpus, field = "featureName") <- NULL`.

In *Step 2*, we begin by selecting the lexicons to include. Herunder, the built-in lexicons and valence word lists are first loaded. The well-known Loughran and McDonald (Loughran and McDonald 2011) and Henry (Henry 2008) word lists together with an English valence word list are supplied to `setup_lexicons()`.⁷ We also add a self-made lexicon of the word ‘uncertainty’ and synonyms, and attach a higher than usual negative sentiment score to these words. We proceed with $L = 3$.

```
R> data("lexicons")
R> data("valence")
R> mine <- data.frame(w = c("uncertainty", "anxiety", "concern",
+                           "distrust", "worries"),
+                   s = c(-2, -2, -2, -2, -2))
R> lexiconsIn <- c(list(myLexicon = mine), lexicons[c("LM_eng", "HENRY_eng")])
R> lex <- setup_lexicons(lexiconsIn = lexiconsIn,
+                       valenceIn = valence[["valence_eng"]])
```

The lexicons look like below. This format of the word lists are `data.tables`, similar to what is used in the `sentimentr` package. The second half of the lexicon represents the original words copied but preceded by a ‘NOT_’. When computing sentiment, all negators in the valence word list (*i.e.*, in `lex[["valence"]]`) are searched for in the corpus texts and also replaced by ‘NOT_’, glued together to the next word. This way, the lexicon can be applied to the modified corpus still considering all words as unigrams. The underlying assumption is that all negators have the same impact, multiplying the lexicon score by -1 . The same logic holds for other valence shifters, such as amplifiers (*e.g.*, very).

```
R> lex[["myLexicon"]]
```

	x	y
1:	anxiety	-2
2:	concern	-2
3:	distrust	-2
4:	uncertainty	-2
5:	worries	-2
6:	NOT_anxiety	2
7:	NOT_concern	2
8:	NOT_distrust	2
9:	NOT_uncertainty	2
10:	NOT_worries	2

The sentiment calculation in *Step 2* and the aggregation in *Step 3* are performed in conjunction with each other. We aggregate sentiment at a monthly frequency, weight the words for within-document aggregation based on the term frequency-inverse document frequency statistic, weight the documents for across-document aggregation proportionally w.r.t. to the number of words in the document, and have an equally-weighted, a linear and several Almon polynomial time aggregation schemes for a lag of 12 months. The last gives $B = 8$. We ignore

⁷As for now, the built-in valence word lists only include negators.

documents with zero sentiment for across-document aggregation, and fill missing dates with a value of zero for across-time aggregation, as per default.

```
R> ctrAgg <- ctr_agg(howWithin = "tf-idf",
+                   howDocs = "proportional",
+                   howTime = c("equal_weight", "linear", "almon"),
+                   do.ignoreZeros = TRUE,
+                   by = "month",
+                   fill = "zero",
+                   lag = 12,
+                   ordersAlm = 1:3,
+                   do.inverseAlm = TRUE,
+                   do.normalizeAlm = TRUE)
```

All is set to construct the sentiment measures, done by calling the `sento_measures()` function. The `summary()` generic displays a brief overview of the composition of the sentiment time series. The output *sentomeasures* object is a classed list with as most important elements "measures" (the textual sentiment time series), "sentiment" (the original sentiment scores per document) and "stats" (a selection of summary statistics). Alternatively, the same output can be obtained with the functions `compute_sentiment()` and `perform_agg()`.

```
R> sentMeas <- sento_measures(corpus, lexicons = lex, ctr = ctrAgg)
R> summary(sentMeas)
```

This *sentomeasures* object contains 144 textual sentiment time series with 226 observations each, at a monthly frequency.

The corpus has following features: wsj wapo economy noneconomy el war

Following lexicons were used to calculate sentiment: myLexicon LM_eng HENRY_eng
 Following scheme was applied for aggregation within documents: tf-idf
 Following scheme was applied for aggregation across documents: proportional
 Following schemes were applied for aggregation across time: equal_weight
 linear almon1 almon1_inv almon2 almon2_inv almon3 almon3_inv

Aggregate statistics:

mean	sd	max	min	meanCorr
-0.007927479	0.005220577	0.003267185	-0.020292885	0.195079920

A *sentomeasures* object can be easily plotted across each of its dimensions. For example, Figure 1 shows a time series of average sentiment for every feature involved.

```
R> p <- plot(sentMeas, group = "features") +
+   guides(colour = guide_legend(nrow = 1))
+ p
```

Step 4 consists of the modelling. One built-in target variable is a news-based index of economic policy uncertainty, the EPU index (Baker *et al.* 2016). We will analyse how well the




Figure 1: Textual sentiment time series averaged across features

sentiment in the texts of our corpus relates to this index. The length of the dependent variable should be equal to the number of observations in the sentiment measures. We omit other explanatory variables but they can easily be added as a named `data.frame` via the `x` argument of `sento_model()`.

```
R> data("epu")
R> y <- epu[epu[["date"]] >= sentMeas[["measures"]][1]$date, ]$index
R> length(y) == nrow(sentMeas[["measures"]])
```

```
[1] TRUE
```

Let us now apply the iterative analysis to illustrate the simplicity in use compared with the rich output. The modelling control is set up below. We made sure the months of `epu` coincide with the months in `sentMeas`, such that `h = 1` means the model aims at forecasting next month's EPU index value. The model is linear and the calibration is based on a Bayesian-like information criterion. We set a sample size of 36 months for a total sample of 226 observations. The `do.parallel` option is set to `TRUE`. The potential `alpha` values all correspond to a good mix between the Ridge and the LASSO regularization.

```
R> ctrIter <- ctr_model(model = "gaussian",
+                       type = "BIC",
+                       h = 1,
+                       alphas = c(0.3, 0.5, 0.7),
+                       do.iter = TRUE,
```

```
+               nSample = 36,
+               do.parallel = TRUE)
```

The **doParallel** package ([Analytics and Weston 2015a](#)) is used to register the backend for parallel computation. The output of the `sent_model()` call is a *sentomodeliter* object, since `do.iter = TRUE`.

```
R> require(doParallel)
R> cl <- makeCluster(detectCores() - 1)
R> registerDoParallel(cl)
R> out <- sent_model(sentMeas, y, ctr = ctrIter)
R> stopCluster(cl)
R> summary(out)
```

Model specifications

- - - - -

Model type: gaussian

Calibration: via BIC information criterion

Sample size: 36

Total number of iterations/predictions: 189

Optimal average elastic net alpha parameter: 0.4777778

Optimal average elastic net lambda parameter: 12.17622

Out-of-sample performance

- - - - -

Mean directional accuracy: 39.3617 %

Root mean squared prediction error: 41.4602

Mean absolute deviation: 29.13337

In-sample performance

- - - - -

Average fraction of deviance explained: 54.75907 %

A more detailed view of the different performance measures, in this case directional accuracy, root mean squared and absolute errors, is obtained via `out[["performance"]]`. A list of the individual *sentomodel* objects is found under `out[["models"]]`. A simple plot to visualize the out-of-sample fit of any *sentomodeliter* object requires only `plot()`. We see in Figure 2 that the generated textual sentiment times series from two U.S. journals only do a decent job in forecasting the EPU index.

```
R> plot(out)
```

The final step, *Step 5*, mainly focuses on the attribution analysis. It can be invoked by the `retrieve_attributions()` function with the modelling object and the related sentiment




Figure 2: Realized EPU index values versus out-of-sample predictions

measures object as inputs. It generates attributions for all out-of-sample dates at the level of individual documents, lexicons, features and time. The function can be applied for `sentomodel` objects as well, for any specific dates using the `refDates` argument. The `do.normalize` argument if set to `TRUE` normalizes all attributions between `-1` and `1`.

```
R> attributions <- retrieve_attributions(out, sentMeas, do.normalize = FALSE)
R> names(attributions)
```

```
[1] "documents" "lexicons" "features" "time"
```

Plotting out the attributions is effortless with the `plot_attributions()` function, according to any of the dimensions, except for individual documents. The attributions are displayed stacked on top of each other, per date. Inspecting Figure 3, one sees for example a clearly evolving influence of the Wall Street Journal feature on the predictions.

```
R> a <- plot_attributions(attributions, group = "features") +
+   guides(fill = guide_legend(nrow = 1))
R> a
```

3.3. Merging functionalities

Apart from several convenience functions such as `predict()`, `scale()`, `fill_measures()` and `select_measures()`, the biggest non-discussed functionality is the further merging of




Figure 3: Prediction attribution to each of the six features

the sentiment measures. For instance, all sentiment measures composed of three particular time aggregation schemes can be merged together as an average. The merging functionality is foremost a way to either diminish or expand the dimensionality of the sentiment measures. It is carried out in two steps: first by specifying another, yet simple, control function, and then by the actual `merge_measures()` call which returns a new `sentomeasures` object. Here, both built-in lexicons, both news sources and two time weighting schemes are respectively merged, or else said collapsed, into one. The `do.keep = FALSE` option means the original measures are not kept after merging, such that the number of sentiment time series effectively goes down.

```
R> ctrMerge <- ctr_merge(sentMeas,
+                         time = list(W = c("equal_weight", "linear")),
+                         lexicons = list(LEX = c("LM_eng", "HENRY_eng")),
+                         features = list(journals = c("wsj", "wapo")),
+                         do.keep = FALSE)
R> sentMeasMerged <- merge_measures(ctrMerge)
R> sentMeasMerged[c("features", "lexicons", "time")]
```

```
$features
[1] "economy" "noneconomy" "el" "war" "journals"
```

```
$lexicons
[1] "myLexicon" "LEX"
```




Figure 4: Globally merged sentiment measure

```
$time
[1] "almon1" "almon1_inv" "almon2" "almon2_inv" "almon3" "almon3_inv" "W"
```

The sentiment measures can also be merged into one measure we refer to as global sentiment. Each of the different components of the dimensions has to receive a weight that stipulates the importance in the global sentiment index. In below example, the highest weight is put on the initial features as well as on the McDonald and Loughran lexicon. The output is a `data.frame`, and can thus be easily used as a variable in any other regression analysis. The global sentiment measure is plotted in Figure 4. The drop in sentiment during the worldwide financial crisis is most apparent.

```
R> glob <- to_global(sentMeas,
+                   lexicons = c(0.10, 0.70, 0.20),
+                   features = c(0.20, 0.20, 0.20, 0.20, 0.10, 0.10),
+                   time = 1)
R> g <- ggplot(data = glob, aes(x = as.Date(row.names(glob)), y = global)) +
+   geom_line() +
+   scale_x_date(name = "Date") +
+   scale_y_continuous(name = "Sentiment") +
+   ggthemes::theme_tufte(base_size = 12) +
+   theme(legend.title = element_blank())
R> g
```


4. Conclusion and future directions

The **sentometrics** package provides a means to calculate sentiment for a large number of texts, aggregate the textual sentiment into a vast quantity of time series at a desired frequency, and use this information in a flexible prediction model. It can thus be deployed to relate a corpus of texts, qualitative data, to a quantitative target variable and retrieve information on which type of sentiment is most informative. This package is a first step towards a consistent econometric analysis in R of the added value of information in texts. In the future, the links with current text mining packages will be enforced to improve the transition from pure textual analysis, corpus construction and feature generation to the econometric analysis. On the last note, we aim to expand the number of available models, as well as estimation and calibration approaches. Lastly, we will expand and improve the textual sentiment calculation engine(s). If you use R or **sentometrics**, please cite the software in publications. In case of the latter, use `citation("sentometrics")`.

Computational details

The example was produced using R 3.4.2 (R Core Team 2017). Code for the illustration is available in the R script `run_vignette.R` located in the `examples` folder on the dedicated GitHub repository.

R itself and all packages used are available from CRAN at <http://CRAN.R-project.org/>. The version under development is available on the **sentometrics** website at <http://sborms.github.io/sentometrics/>, which provides additional examples of how to use the package.

Acknowledgments

We acknowledge Google Summer of Code 2017 and Fonds des Donations at the University of Neuchâtel for their financial support.

References

- Analytics R, Weston S (2015a). *doParallel: Foreach Parallel Adaptor for the ‘parallel’ Package*. R package version 1.0.10, URL <https://CRAN.R-project.org/package=doParallel>.
- Analytics R, Weston S (2015b). *foreach: Provides Foreach Looping Construct for R*. R package version 1.4.3, URL <https://CRAN.R-project.org/package=foreach>.
- Ardia D, Bluteau K, Boudt K (2017). “Questioning the news about economic growth: Sparse forecasting using thousands of news-based sentiment values.” *Working paper*.
- Arnold T (2017). “cleanNLP: A Tidy Data Model for Natural Language Processing using cleanNLP.” R package version 1.10.0, URL <https://CRAN.R-project.org/package=cleanNLP>.

- Baker S, Bloom N, Davis S (2016). “Measuring Economic Policy Uncertainty.” *NBER Working paper*.
- Benoit K, Watanabe K, Nulty P, Obeng A, Wang H, Lauderdale B, Lowe W (2017). *quanteda: Quantitative Analysis of Textual Data*. R package version 0.99.12, URL <http://quanteda.io>.
- Caporin M, Poli F (2017). “Building news measures from textual data and an application to volatility forecasting.” *Econometrics*, **5**(3), 35. doi:10.3390/econometrics5030035.
- Catania L, Bernardi M (2017). *MCS: Model Confidence Set Procedure*. R package version 0.1.3, URL <https://CRAN.R-project.org/package=MCS>.
- Davis A, Piger J, Sedor L (2012). “Beyond the numbers: Measuring the information content of earnings press release language.” *Contemporary Accounting Research*, **29**, 845–868. doi:10.1111/j.1911-3846.2011.01130.x.
- Dowle M, Srinivasan A (2017). *data.table: Extension of ‘data.frame’*. R package version 1.10.4-2, URL <https://CRAN.R-project.org/package=data.table>.
- Feinerer I, Hornik K, Meyer D (2008). “Text Mining Infrastructure in R.” *Journal of Statistical Software*, **22**(5), 1–54. doi:10.18637/jss.v025.i05.
- Friedman J, Hastie T, Tibshirani R (2010). “Regularization Paths for Generalized Linear Models via Coordinate Descent.” *Journal of Statistical Software*, **33**(1), 1–22. doi:10.18637/jss.v033.i01.
- Hansen P, Lunde A, Nason J (2011). “The model confidence set.” *Econometrica*, **79**, 453–497. doi:10.3982/ECTA5771.
- Henry E (2008). “Are investors influenced by how earnings press releases are written?” *Journal of Business Communication*, **45**(4), 363–407. doi:10.1177/0021943608319388.
- Hoerl A, Kennard R (1970). “Ridge regression: Biased estimation for nonorthogonal problemst.” *Technometrics*, **12**, 55–67. doi:10.1080/00401706.1970.10488634.
- Kuhn M (2017). *caret: Classification and Regression Training*. R package version 6.0-77, URL <https://CRAN.R-project.org/package=caret>.
- Loughran T, McDonald B (2011). “When is a liability not a liability? Textual analysis, dictionaries, and 10-Ks.” *Journal of Finance*, **66**(1), 35–65. doi:10.1111/j.1540-6261.2010.01625.x.
- R Core Team (2017). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. R version 3.4.2, URL <https://www.R-project.org/>.
- Rinker TW (2013). *qdap: Quantitative Discourse Analysis Package*. R package version 2.2.5, URL <http://github.com/trinker/qdap>.
- Rinker TW (2017). *sentimentr: Calculate Text Polarity Sentiment*. R package version 1.0.0, URL <http://github.com/trinker/sentimentr>.

- Silge J, Robinson D (2016). “tidytext: Text Mining and Analysis Using Tidy Data Principles in R.” **1**(3). doi:10.21105/joss.00037.
- Tetlock P, Saar-Tsechansky M, Macskassy S (2008). “More than words: Quantifying language to measure firms’ fundamentals.” *Journal of Finance*, **63**, 1437–1467. doi:10.1111/j.1540-6261.2008.01362.x.
- Tibshirani R (1996). “Regression shrinkage and selection via the LASSO.” *Journal of the Royal Statistical Society: Series B*, **58**, 267–288. doi:10.1111/j.1467-9868.2011.00771.x.
- Tibshirani R, Taylor J (2012). “Degrees of freedom in LASSO problems.” *Annals of Statistics*, **4**, 1198–1232. doi:10.1214/12-AOS1003.
- Zou H, Hastie T (2005). “Regularization and variable selection via the elastic net.” *Journal of the Royal Statistical Society: Series B*, **67**, 301–320. doi:10.1111/j.1467-9868.2005.00503.x.

Affiliation:

David Ardia
Institute of Financial Analysis
University of Neuchâtel, Switzerland
& Department of Finance, Insurance and Real Estate
Laval University, Canada
E-mail: david.ardias@unine.ch

Keven Bluteau
Institute of Financial Analysis
University of Neuchâtel, Switzerland
& Vrije Universiteit Brussel, Belgium
E-mail: keven.bluteau@unine.ch

Samuel Borms
Institute of Financial Analysis
University of Neuchâtel, Switzerland
& Vrije Universiteit Brussel, Belgium
E-mail: samuel.borms@unine.ch

Kris Boudt
Solvay Business School
Vrije Universiteit Brussel, Belgium
& Vrije Universiteit Amsterdam, The Netherlands
E-mail: kris.boudt@vub.be